

# **AI-Powered Customer Support System: Complete Architecture & Design**

## System Architecture Overview

The AI-powered customer support system utilizes a multi-agent architecture with specialized components working together to automate and enhance the support workflow. Here's the complete system design from an AI/ML perspective:

Unsupported image

## High-Level Architecture

The system follows a modular design with three main layers:

### 1. **Presentation Layer**

- Web interface (HTML, CSS, JavaScript)
- RESTful API endpoints

### 2. **Business Logic Layer**

- Multi-agent system
- ML models
- Knowledge base integration
- Workflow orchestration

### 3. **Data Layer**

- PostgreSQL database
- Ticket repository
- Knowledge base storage
- User and authentication data

## AI/ML Components in Detail

### 1. Natural Language Processing Core

At the heart of the system is the Natural Language Processing (NLP) capability, which:

- Processes unstructured text from customer tickets
- Extracts key information (sentiment, entities, intent)
- Provides embeddings for similarity matching
- Powers conversational interactions

## Implementation Details:

- **Primary LLM Integration:** Ollama-based on-premise models

- **Fallback Mechanism:** Rule-based responses when LLM is unavailable
- **Embedding Generation:** Vector-based representation for semantic similarity

## 2. Multi-Agent System

The system implements a collaborative multi-agent architecture where each agent has a specialized role:

ClassifierAgent

```
def classify_ticket(self, description):
    """Classify a ticket based on its description"""

    # 1. Extract key features from text
    # 2. Apply ML classification model
    # 3. Determine sentiment analysis
    # 4. Calculate priority based on keywords and sentiment
    # 5. Generate concise summary
    # 6. Return structured classification data
```

**Working:**

1. Receives raw ticket text from frontend
2. Uses a trained TF-IDF vectorizer + classification model
3. Categorizes into predefined issue types (software, hardware, network, account)
4. Extracts priority signals from text
5. Applies sentiment analysis to detect customer frustration
6. Returns structured classification for routing

ResolutionAgent

```
def suggest_solutions(self, ticket):
    """Suggest solutions for a given ticket"""

    # 1. Check knowledge base for similar issues
    # 2. Retrieve historical solutions for similar tickets
    # 3. Generate new solution using LLM if no matches
    # 4. Rank solutions by success probability
    # 5. Return formatted solution suggestions
```

**Working:**

1. Takes classified ticket as input
2. First queries knowledge base using semantic search
3. Retrieves similar past tickets using feature similarity
4. Generates solution suggestions using combined approach
5. Ranks solutions based on historical success rates
6. Returns formatted solution suggestions to frontend

#### EscalationAgent

```
def should_escalate(self, ticket, conversation_history=None):
```

```
    """Determine if a ticket needs human intervention"""
```

- # 1. Analyze conversation complexity
- # 2. Detect unresolved issues after multiple attempts
- # 3. Identify technical jargon or specific product mentions
- # 4. Check for escalation keywords/phrases
- # 5. Return escalation decision with reason

#### Working:

1. Monitors ongoing conversations
2. Detects signals requiring human expertise
3. Identifies patterns indicating customer frustration
4. Recognizes complex technical issues beyond automation
5. Triggers escalation workflow when threshold is crossed

#### FeedbackAgent

```
def process_feedback(self, ticket_id, rating, comment):
```

```
    """Process feedback for a ticket"""
```

- # 1. Record feedback in database
- # 2. Update solution success rates
- # 3. Identify improvement opportunities
- # 4. Create knowledge base entries for successful resolutions
- # 5. Trigger learning for ML components

#### Working:

1. Processes customer feedback on resolutions
2. Updates success metrics for suggested solutions
3. Feeds data back to improve classification models
4. Generates knowledge base entries from successful resolutions
5. Provides continuous learning signal to ML models

#### ChatbotAgent

```
def respond_to_query(self, user_message, conversation_history=None, session_id=None):
```

```
    """Generate a response based on conversation state and user query"""
```

- ```
    # 1. Determine conversation state
    # 2. Process message based on current state
    # 3. Generate contextual response
    # 4. Update conversation state
    # 5. Return response with action metadata
```

#### Working:

1. Maintains conversational state for structured flow
2. Guides users through support workflow
3. Detects intents from natural language
4. Suggests troubleshooting steps for common issues
5. Creates tickets when issues can't be immediately resolved

### 3. Machine Learning Models

#### TicketClassifier

```
def predict_category(self, description):
```

```
    """Predict the category of a ticket based on its description"""
```

- ```
    # 1. Preprocess text data
    # 2. Extract features using TF-IDF
    # 3. Apply trained classifier
    # 4. Return predicted category with confidence
```
- **Algorithm:** TF-IDF vectorization + RandomForest classifier
  - **Features:** N-grams, keyword presence, text length, technical terms

- **Training:** Supervised learning on historical ticket data
- **Performance Metrics:** Accuracy, F1-score, confusion matrix

SentimentAnalyzer

```
def analyze_sentiment(self, text):
```

```
    """Analyze the sentiment of text with enhanced features"""
```

```
    # 1. Preprocess text
```

```
    # 2. Apply sentiment lexicon and rules
```

```
    # 3. Identify emotional signals
```

```
    # 4. Calculate sentiment score
```

```
    # 5. Return structured sentiment data
```

- **Algorithm:** Lexicon-based approach with rule-based augmentation
- **Features:** Emotion words, negations, intensifiers, punctuation patterns
- **Output:** Sentiment category (Positive, Neutral, Negative, Frustrated, Urgent)
- **Applications:** Priority scoring, escalation decisions, conversation health

ResolutionPredictor

```
def predict_resolution(self, description, category):
```

```
    """Predict the resolution for a ticket"""
```

```
    # 1. Find similar historical tickets
```

```
    # 2. Extract common resolution patterns
```

```
    # 3. Generate resolution suggestion
```

```
    # 4. Return formatted resolution text
```

- **Algorithm:** Hybrid of similarity matching and generative AI
- **Features:** Issue description, category, similar historical resolutions
- **Training:** Supervised learning on resolved ticket pairs (problem-solution)
- **Performance Metric:** Resolution acceptance rate, resolution time reduction

ConversationHealthAnalyzer

```
def analyze_conversation(self, conversation_history):
```

```
    """Analyze a conversation and return health metrics"""
```

```
    # 1. Extract turn-taking patterns
```

# 2. Analyze sentiment progression

# 3. Detect loop patterns or repetitions

# 4. Calculate response times

# 5. Return structured health metrics

- **Algorithm:** Rule-based system with statistical pattern recognition
- **Features:** Turn distribution, sentiment change, repetition rate, message length
- **Output:** Health score (0-100) and improvement suggestions
- **Application:** Quality monitoring and agent training

#### 4. Knowledge Base Integration

The knowledge base system serves as the institutional memory with:

```
def create_knowledge_base_entry_from_ticket(ticket_id, admin_id=1):  
    """Create a knowledge base entry from a successfully resolved ticket"""  
  
    # 1. Retrieve ticket and conversation data  
  
    # 2. Format problem description from user messages  
  
    # 3. Extract solution from resolved ticket  
  
    # 4. Generate tags for searchability  
  
    # 5. Create structured KB entry
```

```
def find_knowledge_base_entries_for_issue(description, category=None, limit=3):  
    """Find relevant knowledge base entries for a given issue description"""  
  
    # 1. Extract keywords from description  
  
    # 2. Match against knowledge base entries  
  
    # 3. Score matches based on relevance  
  
    # 4. Return top-matching entries
```

#### Working:

1. Continuously learns from successfully resolved tickets
2. Automatically creates structured knowledge entries
3. Provides semantic search capabilities
4. Feeds solutions back into the resolution process
5. Improves over time with feedback and usage statistics

## Data Flow and Processing Pipeline

### 1. Ticket Creation & Classification Flow

1. User submits issue description via web interface
2. Frontend sends API request to /api/tickets endpoint
3. Backend invokes ClassifierAgent.classify\_ticket()
4. ML model processes text and returns classification
5. System creates ticket with metadata and routes appropriately
6. User receives confirmation with estimated resolution time

### 2. Resolution Suggestion Flow

1. Support agent/user views ticket details
2. System calls ResolutionAgent.suggest\_solutions()
3. Knowledge base is queried for similar issues
4. Historical solutions are retrieved and ranked
5. LLM generates additional suggestions if needed
6. Suggested solutions are displayed to agent/user

### 3. Continuous Learning Flow

1. User provides feedback on resolution
2. FeedbackAgent.process\_feedback() processes rating/comments
3. Solution success rates are updated
4. Positive resolutions create knowledge base entries
5. ML models periodically retrain on expanded dataset
6. System performance improves over time

## Technical Implementation Details

### On-Premise LLM Integration

class OllamaClient:

"""Client for interacting with Ollama API for LLM capabilities"""

def generate(self, prompt, system\_prompt=None):

"""Generate a response using the Ollama API or fallback to rule-based responses"""

- # 1. Prepare request with prompt and system context
- # 2. Send request to Ollama API
- # 3. Process and return response
- # 4. Fall back to rule-based system if unavailable
- **Model Hosting:** On-premise via Ollama
- **Fallback Mechanism:** Rule-based generation when LLM unavailable
- **Prompt Engineering:** Specialized prompts for different agent functions
- **Performance Considerations:** Caching, request batching, response streaming

### Database Schema Design

The database schema supports the ML/AI functionality with:

- **Tickets table:** Stores classification metadata, sentiment analysis
- **Conversations table:** Contains message history for contextual processing
- **Solutions table:** Tracks suggested solutions and success rates
- **KnowledgeBase table:** Stores curated solution entries with semantic tags
- **Feedback table:** Captures user ratings and comments for learning

### API Integration Layer

The API layer connects frontend and AI components:

```
@app.route('/api/tickets/<ticket_id>/suggest-solutions', methods=['GET'])
```

```
def suggest_solutions(ticket_id):
```

```
    """API endpoint to get solution suggestions for a ticket"""
```

```
    # 1. Find ticket in database
```

```
    # 2. Query knowledge base for relevant entries
```

```
    # 3. Get ML-generated solutions
```

```
    # 4. Combine and return responses
```

- **RESTful Endpoints:** Structured API for agent interactions
- **Authentication:** Role-based access control via Flask-Login
- **Error Handling:** Graceful degradation when AI services unavailable
- **Response Formats:** Standardized JSON with metadata and confidence scores

### Performance Optimization



## ML Model Efficiency

- **Lazy Loading:** Models initialize only when needed
- **Feature Selection:** Optimized feature sets for classification
- **Caching:** Frequent queries cached for performance
- **Batch Processing:** Where possible, batch requests to models

## Scalability Considerations

- **Stateless Design:** Agents maintain minimal state between requests
- **Database Indexing:** Optimized for text search and retrieval
- **Knowledge Base Partitioning:** Category-based organization for faster retrieval
- **Asynchronous Processing:** Background tasks for non-critical operations

## Practical System Limitations and Mitigations

1. **LLM Availability:** Fallback to rule-based responses when unavailable
2. **Cold Start Problem:** Preloaded knowledge base with common solutions
3. **Classification Errors:** Confidence thresholds with human verification
4. **Privacy Concerns:** On-premise processing of all sensitive data
5. **Hallucination Risks:** Knowledge base grounding for generated content

## Interview Questions and Answers

1. How does your system handle the cold start problem for ML models?

**Answer:** We address the cold start problem through multiple strategies:

1. **Preloaded Knowledge Base:** We seed the system with high-quality, manually curated knowledge base entries covering common issues.
2. **Initial Model Training:** Our models are pre-trained on a dataset of historical support tickets (from `Historical_ticket_data.csv`) to provide baseline classification performance.
3. **Rule-Based Fallbacks:** When confidence is low or data is insufficient, we fall back to rule-based approaches for classification and resolution.
4. **Active Learning:** We prioritize human review of low-confidence predictions, using these reviews to improve model performance.
5. **Transfer Learning:** We leverage pre-trained language models that already have broad language understanding, then fine-tune on our domain-specific data.

This hybrid approach ensures the system provides value immediately while continuously improving as more data becomes available.

2. What measures have you taken to ensure data privacy and security in your AI implementation?

**Answer:**

1. **On-Premise LLM:** By using Ollama for on-premise LLM hosting, we ensure sensitive customer data never leaves the organization's infrastructure.
2. **Role-Based Access Control:** We implement strict role-based access in our models where customer data is only accessible to authorized personnel.
3. **Data Anonymization:** When using data for training, we anonymize personal identifiers.
4. **Minimized Data Collection:** Our agents are designed to process only the information necessary for their specific tasks.
5. **Secure Knowledge Base:** Our knowledge base contains generalized solutions without customer-specific information.
6. **Audit Logging:** All AI-driven decisions and data accesses are logged for accountability.
7. **User Consent:** Clear disclosure when interacting with automated systems with options to escalate to human support.

3. How do your ML models improve over time, and what metrics do you use to measure their performance?

**Answer:** Our system employs continuous learning through:

1. **Feedback Loop Integration:** User feedback directly influences model training through the FeedbackAgent.
2. **Periodic Retraining:** Models are retrained on expanded datasets as new tickets and resolutions accumulate.
3. **A/B Testing:** New model versions are tested against existing ones before deployment.

For performance measurement, we track multiple metrics:

- **Classification Accuracy:** Percentage of tickets correctly categorized
- **Resolution Success Rate:** Percentage of suggested solutions that resolved issues
- **Time-to-Resolution:** Average time from ticket creation to resolution
- **Escalation Rate:** Percentage of tickets requiring human intervention
- **Customer Satisfaction:** Ratings and sentiment analysis of feedback
- **Knowledge Base Utilization:** Rate at which knowledge base entries successfully resolve issues

We use a weighted combination of these metrics to holistically evaluate system performance rather than optimizing for any single metric.

4. How does your system decide when to escalate a ticket to a human agent?

**Answer:** The EscalationAgent makes this decision using a multi-factor approach:

1. **Confidence Threshold:** If the classification or resolution confidence falls below predetermined thresholds.
2. **Conversation Analysis:** Detecting patterns indicating customer frustration, such as repetition, increasing negative sentiment, or explicit escalation requests.
3. **Issue Complexity:** Identifying technical language or specific product mentions that suggest specialized knowledge requirements.
4. **Repeated Solutions:** If multiple suggested solutions have failed to resolve the issue.
5. **User Status:** VIP customers or those with urgent issues may be escalated earlier.
6. **Business Rules:** Certain issue categories (e.g., security concerns, billing disputes) are always escalated.

The system uses a weighted scoring mechanism for these factors, and when the combined score exceeds the escalation threshold, the ticket is routed to a human agent with context from the automated interaction.

5. What technical challenges did you face in implementing the multi-agent architecture, and how did you overcome them?

**Answer:** Several significant challenges emerged:

1. **Agent Coordination:** Ensuring agents worked together cohesively without conflicting decisions.  
*Solution:* Implemented a well-defined workflow with clear handoffs between agents and a central orchestration mechanism.
2. **State Management:** Maintaining conversation context across multiple interactions.  
*Solution:* Designed a session-based state management system in the ChatbotAgent.
3. **Error Propagation:** Preventing cascading failures when one agent produced incorrect output.  
*Solution:* Added validation layers between agent interactions and implemented fallback mechanisms.
4. **LLM Reliability:** Dealing with occasional unavailability of the Ollama service.  
*Solution:* Created rule-based fallback systems for critical functionality.
5. **Performance Bottlenecks:** Some agent operations were computationally expensive.  
*Solution:* Implemented caching strategies and asynchronous processing for non-urgent tasks.
6. **Continuous Learning Integration:** Balancing model stability with improvement.  
*Solution:* Established a staged deployment process with validation gates before production deployment.

6. How does your knowledge base differ from a simple FAQ system?

**Answer:** Our knowledge base is far more sophisticated than a traditional FAQ system:

1. **Dynamic Growth:** Unlike static FAQs, our knowledge base automatically expands from successful ticket resolutions.
2. **Semantic Search:** We use semantic understanding rather than keyword matching, allowing for concept-based retrieval.
3. **Context-Aware Retrieval:** The system considers ticket category, priority, and user history when retrieving solutions.
4. **Solution Ranking:** Entries are ranked by historical success rate and helpfulness ratings.
5. **Structured Problem-Solution Format:** Each entry contains formatted problem descriptions, solutions, and metadata.
6. **Continuous Improvement:** Entries evolve based on user feedback and success metrics.
7. **ML Integration:** The knowledge base is tightly integrated with our ML models for bidirectional enhancement.
8. **Tag-Based Organization:** Automatic tagging facilitates organization and improves retrievability.

This creates an institutional memory that continuously improves as more support interactions occur, unlike static FAQ systems.

7. What approach did you take to handle the diversity of customer issues in your classification model?

**Answer:** Handling diverse customer issues required a multi-faceted approach:

1. **Hierarchical Classification:** We implemented a two-level classification system - primary category (software, hardware, network, account) followed by subcategory classification.
2. **Feature Engineering:** We extract domain-specific features like error codes, product mentions, and technical terminology.
3. **Text Preprocessing:** Custom preprocessing pipeline that preserves technical terms while normalizing variations.
4. **Ensemble Modeling:** Combining multiple classification algorithms to improve robustness across different issue types.
5. **Confidence Scoring:** Each classification includes a confidence score to identify uncertain cases.
6. **Regular Taxonomy Updates:** The category system evolves as new issue types emerge.
7. **Zero-Shot Learning:** For rare issues, we leverage LLM capabilities to classify without explicit training examples.
8. **Data Augmentation:** For underrepresented categories, we use techniques to synthetically increase training examples.

This approach ensures high accuracy across common issues while maintaining reasonable performance for novel or rare problems.

8. How do you evaluate the quality of AI-generated solutions before presenting them to users?

**Answer:** We employ a multi-stage quality assurance process:

1. **Knowledge Base Grounding:** Solutions are primarily sourced from verified knowledge base entries.
2. **Confidence Scoring:** Each generated solution includes a confidence metric.
3. **Success Rate Tracking:** Historical success rates influence solution ranking.
4. **Content Validation:** Rule-based checks verify that solutions contain actionable steps.
5. **Sensitivity Screening:** Filter out potentially harmful or inappropriate content.
6. **Human Review:** Low-confidence or novel solutions may be flagged for human review.
7. **A/B Testing:** New solution generation strategies are tested with limited audience before full deployment.
8. **User Feedback Loop:** Direct feedback on solution quality drives continuous improvement.

This layered approach ensures that users receive high-quality, relevant solutions while maintaining appropriate human oversight of the AI system.

9. What architectural decisions did you make to ensure the system remains performant as it scales?

**Answer:** We designed for scalability from the ground up:

1. **Modular Agent Design:** Each agent operates independently, allowing for horizontal scaling.
2. **Database Optimization:** Proper indexing, especially for text fields and commonly queried attributes.
3. **Caching Strategy:** Multi-level caching for frequently accessed knowledge base entries and classification results.
4. **Asynchronous Processing:** Non-blocking operations for time-consuming tasks like ticket classification.
5. **Batch Processing:** Where possible, batch requests to ML models to maximize throughput.
6. **Query Optimization:** Careful SQL query design to minimize database load.
7. **Resource Isolation:** Critical paths (customer-facing) are isolated from analytical processing.
8. **Stateless Design:** Agents maintain minimal state, allowing for load balancing across instances.
9. **Scheduled Maintenance:** Background tasks like model retraining occur during off-peak hours.
10. **Performance Monitoring:** Comprehensive metrics to identify bottlenecks early.

These decisions ensure the system maintains responsiveness under increasing load, with clear paths for component-level scaling as needed.

10. How does your system balance automation with maintaining a human-like customer experience?

**Answer:** Achieving this balance was a key design consideration:

1. **Conversation Design:** Our ChatbotAgent uses natural language patterns and proper conversation flow.
2. **Transparency:** We clearly identify automated interactions while maintaining a personable tone.
3. **Seamless Handoff:** When escalation occurs, full context is provided to human agents.
4. **Emotional Intelligence:** SentimentAnalyzer detects customer emotions and adjusts response tone.
5. **Proactive Escalation:** The system recognizes when automation isn't meeting customer needs.
6. **Personalization:** Responses incorporate customer history and preferences.
7. **Continuous Improvement:** Customer satisfaction ratings directly inform our training process.
8. **Human-in-the-Loop:** Critical decisions always have human oversight options.
9. **Response Diversity:** Avoiding repetitive, template-like responses that feel robotic.
10. **Appropriate Limitations:** The system acknowledges its limitations rather than providing incorrect information.

This approach creates a balanced experience where automation handles routine issues efficiently while preserving the human touch for complex situations.